

# VERY DEEP CONVOLUTION NETWORK FOR SUPER RESOLUTION

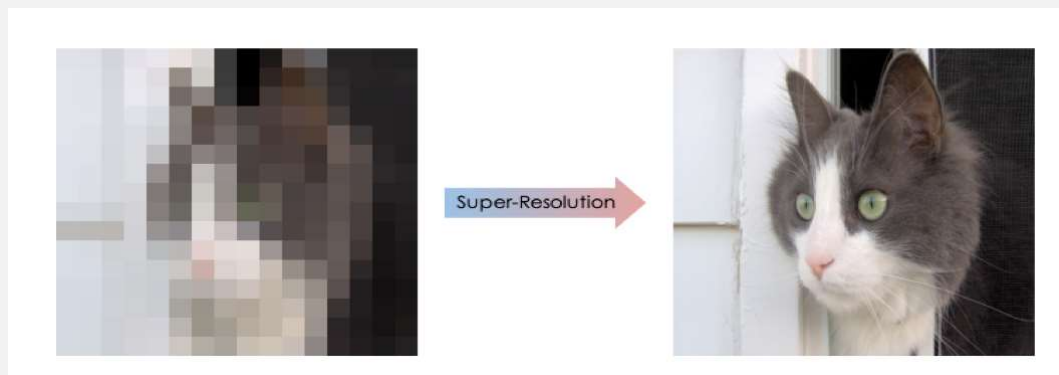


Deep-Learning Project CS 577

# WHAT IS SUPER RESOLUTION

- Before getting into the nitty gritty of super resolutions, let us understand some concepts regarding images.

- What is an image?
- What is resolution?






# INTRODUCTION


Image Super-Resolution: Low Resolution image to High Resolution image

Usefulness: medical field, security, Biological field.

Existing deep learning technique and its problem: SRCNN -small image regions, slow, single scale

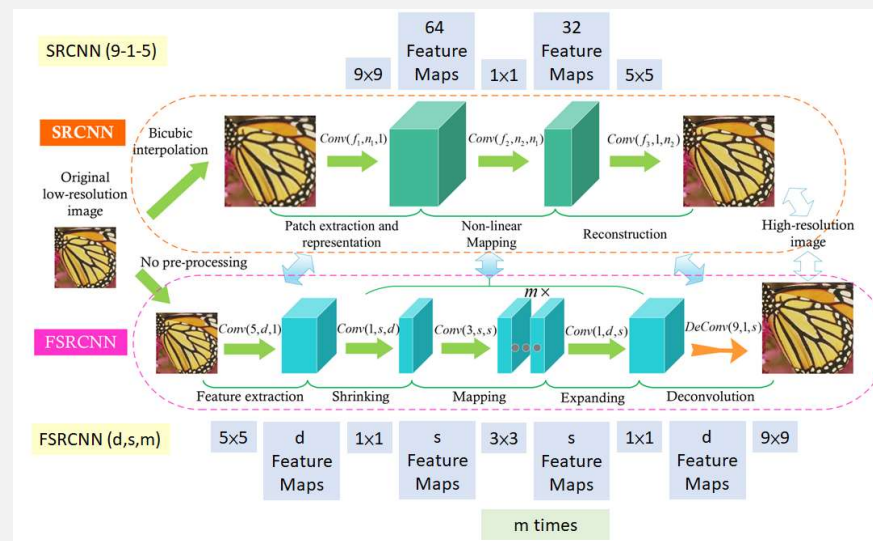


LR image



HR image

# PREVIOUS MODELS



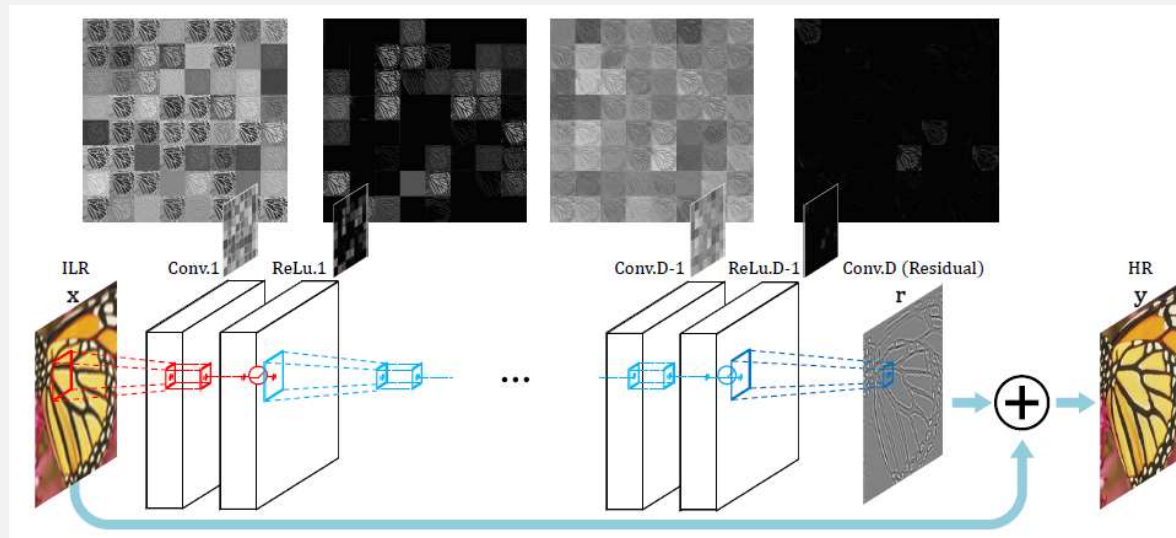
## PROBLEM STATEMENT

How to implement a better Image Super-Resolution network that works better than SRCNN?


SRCNN	Our network
Small Images	Big Images
Slow	Fast
Uni-Scale	Multi-Scale

# VDSR NETWORK

- Use a very deep convolutional network inspired by K.Simonyan and A.Zisserman
- Deep Convolution-layers network



- Issue with using a very deep network to predict dense outputs:
  - Size of the feature map gets reduced every time convolution operations are applied
- Proposed solution:
  - Use Zero Padding before convolutions





# IMPLEMENTATION

```

train_label = train_x #Storing the original resolution img as train label
test_label = test_x #Storing the original resolution img as test label

print("Training set ground truth has shape: " + str(train_label.shape))
print("Test set ground truth has shape: " + str(test_label.shape))

#setting output width x height
w,h = train_label.shape[1:-1]
input_dim = train_label.shape[1:]
scale_factor = scale_f
#size_output_img = (w,h)
print('Width and Height of training and test data should be:',input_dim)
dw= int(w/scale_factor)
dh=int(h/scale_factor)
#print('Scale factor of for width and height:',dw,dh)
ds =(dw,dh)

#Creating Train_X and Test_X from the ground truth
#this is done so by reducing the size of copy of ground truth(train_label, test_label)
#Here the scale is 2, the training shape = 16x16x3 i.e half of 32x32x3
train_data = []
for img in train_label:
    img_re = cv2.resize(img, ds, interpolation=cv2.INTER_CUBIC)
    train_data.append(img_re)
train_data = np.asarray(train_data)

test_data = []
for img in test_label:
    img_re = cv2.resize(img, ds, interpolation=cv2.INTER_CUBIC)
    test_data.append(img_re)
#train_data = np.array(train_data, dtype=float)
b1 = train_data[:10000]
b2 = train_data[10000:20000]
b3 = train_data[20000:30000]
b4 = train_data[30000:40000]
b5 = train_data[40000:50000]
b1 = tf.convert_to_tensor(b1, dtype=float)
b1 = resize_images(b1, scale_factor, scale_factor, 'channels_last', interpolation='bilinear')
b2 = tf.convert_to_tensor(b2, dtype=float)
b2 = resize_images(b2, scale_factor, scale_factor, 'channels_last', interpolation='bilinear')
b3 = tf.convert_to_tensor(b3, dtype=float)
b3 = resize_images(b3, scale_factor, scale_factor, 'channels_last', interpolation='bilinear')
b4 = tf.convert_to_tensor(b4, dtype=float)
b4 = resize_images(b4, scale_factor, scale_factor, 'channels_last', interpolation='bilinear')
b5 = tf.convert_to_tensor(b5, dtype=float)
b5 = resize_images(b5, scale_factor, scale_factor, 'channels_last', interpolation='bilinear')
x_scaled_train= tf.concat([b1,b2,b3,b4,b5],0)
test_data = tf.convert_to_tensor(test_data, dtype=float)

```

X- input image[Low Resolution ]: For the model



Y - Label output[High Resolution :]



After upscaling the input images:  
 Train data shape: (40000, 32, 32, 3)  
 Train label shape: (40000, 32, 32, 3)  
 Vali data shape: (10000, 32, 32, 3)  
 Vali label shape: (10000, 32, 32, 3)  
 Test data shape: (10000, 32, 32, 3)  
 Test label shape: (10000, 32, 32, 3)

```

1 def vdsr(input_dim, l):
2     LR = Input(shape=input_dim, name='input')
3     X = ZeroPadding2D()(LR)
4     X = Conv2D(64, (3,3), name='CONV1')(X)
5     X = ReLU()(X)
6
7     for i in range(l-2):
8         X = ZeroPadding2D()(X)
9         X = Conv2D(64, (3,3), name='CONV%i' % (i+2))(X)
10        X = ReLU()(X)
11
12    X = ZeroPadding2D()(X)
13    residual = Conv2D(1, (3,3), name='CONV%i' % l)(X)
14
15    out = Add()([LR, residual])
16
17    return Model(LR, out)
18
19

```

```

Model: "model"

```

Layer (type)	Output Shape	Param #	Connected to
input (InputLayer)	[(None, 32, 32, 3)]	0	[]
zero_padding2d (ZeroPadding2D)	(None, 34, 34, 3)	0	['input[0][0]']
CONV1 (Conv2D)	(None, 32, 32, 64)	1792	['zero_padding2d[0][0]']
re_lu (ReLU)	(None, 32, 32, 64)	0	['CONV1[0][0]']
zero_padding2d_1 (ZeroPadding2D)	(None, 34, 34, 64)	0	['re_lu[0][0]']
CONV2 (Conv2D)	(None, 32, 32, 64)	36928	['zero_padding2d_1[0][0]']
re_lu_1 (ReLU)	(None, 32, 32, 64)	0	['CONV2[0][0]']
zero_padding2d_2 (ZeroPadding2D)	(None, 34, 34, 64)	0	['re_lu_1[0][0]']
CONV3 (Conv2D)	(None, 32, 32, 64)	36928	['zero_padding2d_2[0][0]']
re_lu_2 (ReLU)	(None, 32, 32, 64)	0	['CONV3[0][0]']
zero_padding2d_3 (ZeroPadding2D)	(None, 34, 34, 64)	0	['re_lu_2[0][0]']
CONV4 (Conv2D)	(None, 32, 32, 64)	36928	['zero_padding2d_3[0][0]']
re_lu_3 (ReLU)	(None, 32, 32, 64)	0	['CONV4[0][0]']
zero_padding2d_4 (ZeroPadding2D)	(None, 34, 34, 64)	0	['re_lu_3[0][0]']
CONV5 (Conv2D)	(None, 32, 32, 1)	577	['zero_padding2d_4[0][0]']
add (Add)	(None, 32, 32, 3)	0	['input[0][0]', 'CONV5[0][0]']

```

=====
Total params: 113,153
Trainable params: 113,153
Non-trainable params: 0

```

# 10, 15 and 20 Layer Deep Models

```
1 model_110 = VDSR10(inp)
2 model_110.summary()
```

Model: "model\_1"

Layer (type)	Output Shape	Param #	Connected to
input_2 (InputLayer)	[(None, 256, 256, 3) 0		
conv2d_5 (Conv2D)	(None, 256, 256, 64) 1792		input_2[0][0]
re_lu_4 (ReLU)	(None, 256, 256, 64) 0		conv2d_5[0][0]
conv2d_6 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_4[0][0]
re_lu_5 (ReLU)	(None, 256, 256, 64) 0		conv2d_6[0][0]
conv2d_7 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_5[0][0]
re_lu_6 (ReLU)	(None, 256, 256, 64) 0		conv2d_7[0][0]
conv2d_8 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_6[0][0]
re_lu_7 (ReLU)	(None, 256, 256, 64) 0		conv2d_8[0][0]
conv2d_9 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_7[0][0]
re_lu_8 (ReLU)	(None, 256, 256, 64) 0		conv2d_9[0][0]
conv2d_10 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_8[0][0]
re_lu_9 (ReLU)	(None, 256, 256, 64) 0		conv2d_10[0][0]
conv2d_11 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_9[0][0]
re_lu_10 (ReLU)	(None, 256, 256, 64) 0		conv2d_11[0][0]
conv2d_12 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_10[0][0]
re_lu_11 (ReLU)	(None, 256, 256, 64) 0		conv2d_12[0][0]
conv2d_13 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_11[0][0]
re_lu_12 (ReLU)	(None, 256, 256, 64) 0		conv2d_13[0][0]
conv2d_14 (Conv2D)	(None, 256, 256, 1) 577		re_lu_12[0][0]
add_1 (Add)	(None, 256, 256, 3) 0		input_2[0][0] conv2d_14[0][0]

\*\*\*\*\*  
Total params: 297,793  
Trainable params: 297,793  
Non-trainable params: 0

```
3 model_115.compile(optimizer=Adam(learning_rate=0.0007,beta_1=0.9), loss='mse', m
```

Model: "model\_2"

Layer (type)	Output Shape	Param #	Connected to
input_3 (InputLayer)	[(None, 256, 256, 3) 0		
conv2d_15 (Conv2D)	(None, 256, 256, 64) 1792		input_3[0][0]
re_lu_13 (ReLU)	(None, 256, 256, 64) 0		conv2d_15[0][0]
conv2d_16 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_13[0][0]
re_lu_14 (ReLU)	(None, 256, 256, 64) 0		conv2d_16[0][0]
conv2d_17 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_14[0][0]
re_lu_15 (ReLU)	(None, 256, 256, 64) 0		conv2d_17[0][0]
conv2d_18 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_15[0][0]
re_lu_16 (ReLU)	(None, 256, 256, 64) 0		conv2d_18[0][0]
conv2d_19 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_16[0][0]
re_lu_17 (ReLU)	(None, 256, 256, 64) 0		conv2d_19[0][0]
conv2d_20 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_17[0][0]
re_lu_18 (ReLU)	(None, 256, 256, 64) 0		conv2d_20[0][0]
conv2d_21 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_18[0][0]
re_lu_19 (ReLU)	(None, 256, 256, 64) 0		conv2d_21[0][0]
conv2d_22 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_19[0][0]
re_lu_20 (ReLU)	(None, 256, 256, 64) 0		conv2d_22[0][0]
conv2d_23 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_20[0][0]
re_lu_21 (ReLU)	(None, 256, 256, 64) 0		conv2d_23[0][0]
conv2d_24 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_21[0][0]
re_lu_22 (ReLU)	(None, 256, 256, 64) 0		conv2d_24[0][0]
conv2d_25 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_22[0][0]
re_lu_23 (ReLU)	(None, 256, 256, 64) 0		conv2d_25[0][0]
conv2d_26 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_23[0][0]
re_lu_24 (ReLU)	(None, 256, 256, 64) 0		conv2d_26[0][0]
conv2d_27 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_24[0][0]
re_lu_25 (ReLU)	(None, 256, 256, 64) 0		conv2d_27[0][0]
conv2d_28 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_25[0][0]
re_lu_26 (ReLU)	(None, 256, 256, 64) 0		conv2d_28[0][0]
conv2d_29 (Conv2D)	(None, 256, 256, 1) 577		re_lu_26[0][0]
add_2 (Add)	(None, 256, 256, 3) 0		input_3[0][0] conv2d_29[0][0]

\*\*\*\*\*  
Total params: 482,433  
Trainable params: 482,433  
Non-trainable params: 0

```
1 model_120 = VDSR20(inp)
2 model_120.summary()
3 model_120.compile(optimizer=Adam(learning_rate=0.0009,beta_1=0.9), loss='mse', metrics=['accu
```

Model: "model\_3"

Layer (type)	Output Shape	Param #	Connected to
input_4 (InputLayer)	[(None, 256, 256, 3) 0		
conv2d_30 (Conv2D)	(None, 256, 256, 64) 1792		input_4[0][0]
re_lu_27 (ReLU)	(None, 256, 256, 64) 0		conv2d_30[0][0]
conv2d_31 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_27[0][0]
re_lu_28 (ReLU)	(None, 256, 256, 64) 0		conv2d_31[0][0]
conv2d_32 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_28[0][0]
re_lu_29 (ReLU)	(None, 256, 256, 64) 0		conv2d_32[0][0]
conv2d_33 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_29[0][0]
re_lu_30 (ReLU)	(None, 256, 256, 64) 0		conv2d_33[0][0]
conv2d_34 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_30[0][0]
re_lu_31 (ReLU)	(None, 256, 256, 64) 0		conv2d_34[0][0]
conv2d_35 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_31[0][0]
re_lu_32 (ReLU)	(None, 256, 256, 64) 0		conv2d_35[0][0]
conv2d_36 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_32[0][0]
re_lu_33 (ReLU)	(None, 256, 256, 64) 0		conv2d_36[0][0]
conv2d_37 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_33[0][0]
re_lu_34 (ReLU)	(None, 256, 256, 64) 0		conv2d_37[0][0]
conv2d_38 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_34[0][0]
re_lu_35 (ReLU)	(None, 256, 256, 64) 0		conv2d_38[0][0]
conv2d_39 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_35[0][0]
re_lu_36 (ReLU)	(None, 256, 256, 64) 0		conv2d_39[0][0]
conv2d_40 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_36[0][0]
re_lu_37 (ReLU)	(None, 256, 256, 64) 0		conv2d_40[0][0]
conv2d_41 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_37[0][0]
re_lu_38 (ReLU)	(None, 256, 256, 64) 0		conv2d_41[0][0]
conv2d_42 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_38[0][0]
re_lu_39 (ReLU)	(None, 256, 256, 64) 0		conv2d_42[0][0]
conv2d_43 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_39[0][0]
re_lu_40 (ReLU)	(None, 256, 256, 64) 0		conv2d_43[0][0]
conv2d_44 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_40[0][0]
re_lu_41 (ReLU)	(None, 256, 256, 64) 0		conv2d_44[0][0]
conv2d_45 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_41[0][0]
re_lu_42 (ReLU)	(None, 256, 256, 64) 0		conv2d_45[0][0]
conv2d_46 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_42[0][0]
re_lu_43 (ReLU)	(None, 256, 256, 64) 0		conv2d_46[0][0]
conv2d_47 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_43[0][0]
re_lu_44 (ReLU)	(None, 256, 256, 64) 0		conv2d_47[0][0]
conv2d_48 (Conv2D)	(None, 256, 256, 64) 36928		re_lu_44[0][0]
re_lu_45 (ReLU)	(None, 256, 256, 64) 0		conv2d_48[0][0]
conv2d_49 (Conv2D)	(None, 256, 256, 1) 577		re_lu_45[0][0]
add_3 (Add)	(None, 256, 256, 3) 0		input_4[0][0] conv2d_49[0][0]

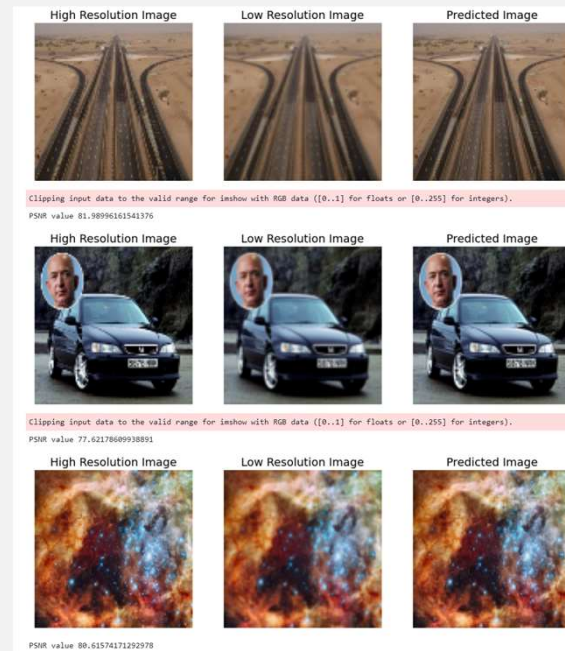
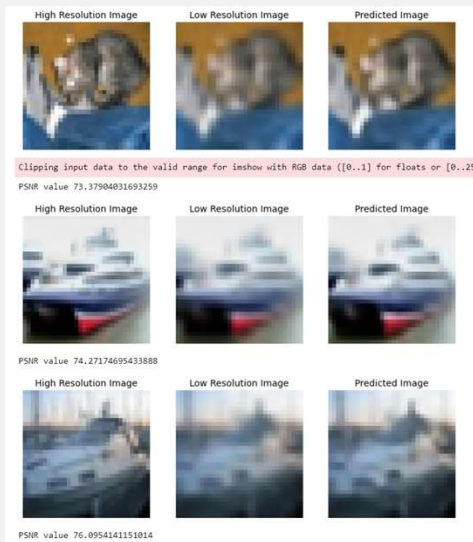
\*\*\*\*\*  
Total params: 667,073  
Trainable params: 667,073  
Non-trainable params: 0

# LOSS

- PSNR : Peak Signal to Noise Ratio

$$\begin{aligned} PSNR &= 10 \cdot \log_{10} \left( \frac{MAX_I^2}{MSE} \right) \\ &= 20 \cdot \log_{10} \left( \frac{MAX_I}{\sqrt{MSE}} \right) \\ &= 20 \cdot \log_{10}(MAX_I) - 10 \cdot \log_{10}(MSE). \end{aligned}$$

# RESULTS AND DISCUSSION





# GUESS THE MODEL?

Maximum psnr value: 94.41647871441019  
Minimum PSNR value: 78.02908746185197  
Average PSNR value: 82.24007084432304

High Resolution Image



Low Resolution Image



Predicted Image



Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).  
PSNR value 82.08008428346795

High Resolution Image



Low Resolution Image



Predicted Image




Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).  
PSNR value 78.02908746185197


High Resolution Image



Low Resolution Image



Predicted Image



PSNR value 80.73586079821227

Maximum psnr value: 92.94281927548559  
Minimum PSNR value: 75.89938128282526  
Average PSNR value: 80.0588363106026

High Resolution Image



Low Resolution Image



Predicted Image



Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).  
PSNR value 79.84382240689713

High Resolution Image



Low Resolution Image




Predicted Image




Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).  
PSNR value 76.3794976067739

High Resolution Image



Low Resolution Image



Predicted Image



PSNR value 77.88437684189171

Maximum psnr value: 94.41647871441019  
 Minimum PSNR value: 78.02908746185197  
 Average PSNR value: 82.24007084432304



Clipping input data to the valid range for imshow with RGB data ([[0..1] for floats or [0..255] for integers) using clip or np.clip.  
 PSNR value 82.08008428346795



Clipping input data to the valid range for imshow with RGB data ([[0..1] for floats or [0..255] for integers) using clip or np.clip.  
 PSNR value 78.02908746185197



PSNR value 80.73586079821227

20 layer Model trained on 100x100x3 images

Maximum psnr value: 92.94281927548559  
 Minimum PSNR value: 75.89938128282526  
 Average PSNR value: 80.0588363106026



Clipping input data to the valid range for imshow with RGB data ([[0..1] for floats or [0..255] for integers) using clip or np.clip.  
 PSNR value 79.84382240689713



Clipping input data to the valid range for imshow with RGB data ([[0..1] for floats or [0..255] for integers) using clip or np.clip.  
 PSNR value 76.3794976067739



PSNR value 77.88437684189171

20 layer Model trained on 256x256x3 images

High Resolution Image



Low Resolution Image



Predicted Image



PSNR value 81.77888095616224

High Resolution Image



Low Resolution Image



Predicted Image



PSNR value 81.9099261310472



Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

High Resolution Image



Low Resolution Image



Predicted Image



Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

PSNR value 81.77888095616224

High Resolution Image



Low Resolution Image



Predicted Image



Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

PSNR value 80.82844640657788

# CONCLUSION

1. After a certain depth, the predicted image will reach its maximum attainable PSNR score.
2. An image of size 32x32x3 or 100x100x3 would not be needing all the 20 layers to create a higher resolution image. By our observing the model's performance we can say, a network of 5 or 10 layers deep is sufficient for small size images.
3. Images which are larger in size would benefit from having deeper networks.
4. Training deep networks of depth = 20 or more will be easier to train for smaller images but would require a lot of computation resources for a larger sized image.
5. We can see the predicted image from 20-layer network, it has sharper edges when compared to the other ones. As we look at bottom edge (near the stomach) region of the bird, we can clearly notice the difference in sharpness produced by the different models. From this we can conclude that lower depth networks provide with a decent PSNR score for small to moderately sized images but to predict finer details for large images, deep networks is a must.

# REFERENCES

- [https://arxiv.org/pdf/1511.04587.pdf?fbclid=IwAR0gm942jqx5Beq9j3maZ90xeEw4qnWNzToiBUKHfCS7OYgTG\\_\\_\\_-hl\\_rx24](https://arxiv.org/pdf/1511.04587.pdf?fbclid=IwAR0gm942jqx5Beq9j3maZ90xeEw4qnWNzToiBUKHfCS7OYgTG___-hl_rx24)
- <http://mmlab.ie.cuhk.edu.hk/projects/SRCNN.html>
- [https://en.wikipedia.org/wiki/Super-resolution\\_imaging](https://en.wikipedia.org/wiki/Super-resolution_imaging)
- [https://www.google.com/url?sa=i&url=https%3A%2F%2Fen.wikipedia.org%2Fwiki%2FPixel\\_aspect\\_ratio&psig=AOvVaw0ohHSfNwsqZYBYPaQEwhaW&ust=1668734123757000&source=images&cd=vfe&ved=0CA8QjRxqFwoTCOiHwomGtPsCFQAAAAAdAAAAABAJ](https://www.google.com/url?sa=i&url=https%3A%2F%2Fen.wikipedia.org%2Fwiki%2FPixel_aspect_ratio&psig=AOvVaw0ohHSfNwsqZYBYPaQEwhaW&ust=1668734123757000&source=images&cd=vfe&ved=0CA8QjRxqFwoTCOiHwomGtPsCFQAAAAAdAAAAABAJ)
- <https://cdn.mos.cms.futurecdn.net/PzPq6Pbn5RqgrWunhEx6rg.jpg>
- [https://miro.medium.com/max/1090/1\\*auit3UsBTjbnzZoLDEJUUg.png](https://miro.medium.com/max/1090/1*auit3UsBTjbnzZoLDEJUUg.png)
- <https://www.scientiamobile.com/wp-content/uploads/2018/12/Pixel-Density.jpg>
- [https://miro.medium.com/max/1060/1\\*rmQjKd4uETJtwYMRTDzmrg.png](https://miro.medium.com/max/1060/1*rmQjKd4uETJtwYMRTDzmrg.png)
- <https://www.google.com/url?sa=i&url=https%3A%2F%2Fzentralwerkstatt.org%2Fblog%2Fpulse&psig=AOvVaw09cargbTBiNkZgNnrNz0DO&ust=1668735205473000&source=images&cd=vfe&ved=0CA8QjRxqFwoTCJD445qJtPsCFQAAAAAdAAAAABAK>
- <https://encrypted-tbn0.gstatic.com/images?q=tbn:ANd9GcRePQoj9vCAU-MxGdkQUWxAIVASBjLWZsFuFewuKNbSnwwEaKpuRV5MM7TtwlpGIPiRnWo&usqp=CAU>

THANK YOU